

Chapter 3: Fault Tolerance

Overview

- What can go wrong
- How to fix it

What can go wrong


- hardware
- software
- environment
- people

⇒ wealth of data in textbook

TANDEM data [turn to textbook page 104, Figure 3.7]

- FT system; many hardware faults masked (only system outages reported)
- environment, operations outages underreported
- MTTF(89) = 20 years (1000 system years, 50 outages)
- MTTF(hardware) = 100 years (10 outages)

MTBF (mean time between failures) informal




Frequency = $\frac{f}{I}$, where f = # failures, I = time

Note: Book calls this *probability* of failure.

Mean Time Between Failures = $MTBF = \frac{I}{f}$

What about MTTF? (Mean Time To Failure)



⇒ if fault distribution is *memoryless*:

$$MTTF = MTBF = \frac{1}{freq} = \frac{I}{f}$$

Example:

two types of faults: red and blue
in 1000 years: 20 red, 30 blue

⇒ $freq_r = \frac{20}{1000}$, $freq_b = \frac{30}{1000}$, $freq_t = \frac{50}{1000}$ where t: either r or b fails.

⇒ $freq_t = freq_r + freq_b$ and $MTTF_t = \frac{1}{freq_r + freq_b}$

Example:

- memory system has 5 boards
- $MTTF_{board} = 10$ years
- system fails if any board fails
- what is the MTTF of the system?

$$MTTF_{system} = \frac{1}{\frac{1}{10} + \frac{1}{10} + \frac{1}{10} + \frac{1}{10} + \frac{1}{10}} \text{ years} = 2 \text{ years}$$


$$\Rightarrow \boxed{MTTF_{N \text{ equal things}} = \frac{MTTF_{each}}{N}}$$

[turn to textbook pages 104 and 106, Figures 3.8, 3.9]

Conclusions:

- need to eliminate high failure rates early & late in component life
- need to consider repairs

Repair Model




$$availability = \frac{MTTF}{MTTF + MTTR}$$

Comments to Figure 3.10 (see textbook page 107)



- disk reliability improving dramatically
1985: MTTF = 8000 hours
1990: MTTF=100,000 hours
- user expectations also increase!
Also: number of disks growing (eg. RAID, large corporations) \Rightarrow still have to take precautions
- should we protect against all failure types?
eg: soft read error? Miss-corrected read?

How to fix it

Before we fix, we need to model failures:



2 approaches:




Options

	Mask	Failsafe&Recovery
Columbus		
Apollo		

⇒ Did you understand MTTF and availability computations for nplex and failstop?

Example 1:

- 2 components A, B with MTTF = 10 years
- no repair
- if one fails, other continues




⇒ 15 years

Example 2

- 3 components, MTTF = 10 years
- TMR (Triple Module Redundancy) = system up until 2 fail
- voter does not fail

(a) no repairs




$$MTTF_{system} = \frac{5}{6} MTTF = 8.3 \text{ years}$$

⇒ worse than 1 component!

⇒ why?

⇒ still useful? (see page 111-112)

(b) with repairs, MTTR = 1 day = 0.003 years



that two components are down happens with probability:

$$(1 - \text{avail})^2 = \left[1 - \frac{MTTF}{MTTF + MTTR} \right]^2 = \left[\frac{MTTR}{MTTF + MTTR} \right]^2 \approx \left[\frac{MTTR}{MTTF} \right]^2$$

$$\Rightarrow \text{freq}_{\text{system}} = \text{freq}_{\text{system-component}} \left[\frac{MTTR}{MTTF} \right]^2 = 3 \text{freq}_{\text{component}} \left[\frac{MTTR}{MTTF} \right]^2$$

$$\Rightarrow MTTF_{\text{system}} = \frac{MTTF}{3} \left[\frac{MTTF}{MTTR} \right]^2 = \frac{10}{3} \left[\frac{10}{0.003} \right]^2 \text{ years} = 3.7 \cdot 10^7 \text{ years}$$


\Rightarrow Real **BIG** Improvement!

Another Problem

- we have 4 components: A, B, C and D; MTTF=1 year; no repairs
- system is fail fast; voters do not fail

Option I


$$MTTF_S = \left(\frac{1}{4} + \frac{1}{3} + \frac{1}{2} + 1 \right) \text{ years}$$
$$= \frac{25}{12} \text{ years} = 2.08 \text{ years}$$



Option II


$$MTTF_X = MTTF_Y = \left(\frac{1}{2} + 1 \right) y = 1.5 y$$

$$MTTF_S = \left(\frac{1.5}{2} + 1.5 \right) y = 2.25 y$$



⇒ should be the same (system runs until all components down)
Problem: A,B,C,D are memoryless (components) but X,Y are not!

N-Plexing



Fail-Vote:

- majority of all components fail \Rightarrow system fails

Fail-Fast:

- all components fail \Rightarrow system fails
- assumes that voter can detect and ignore failed components

Summary so far

- What can go wrong
- How to evaluate reliability (MTTF, Availability)
- Overview of techniques for improving reliability


Next: How to make *storage* reliable.

Building highly available storage (“How to fix it!”)

Overview

- Define model
- Make storage reliable
- Make more reliable

Disk Model



PAGE(i) : contents
STATUS(i) : good or bad (error code)

GET

$(status, block) = GET(address)$

GET(*i*) events

Desired:

IF status(*i*) = good THEN returns (good, page(*i*))
IF status(*i*) = bad THEN returns (bad, -)

Undesired:

Expected:

- status(*i*) = good BUT get(*i*) returns (bad, -)
occurs at most *k* times in a row (soft read error)

Unexpected:

- status(*i*) = bad BUT get(*i*) returns (good, *)
- status(*i*) = good BUT returns (good, page(*j*))
where $j \neq i$ (undetected errors)
- get(*i*) never returns
- get(*i*) causes fire, data center burns up
- ...

Note:

Good Thing:

don't have to worry about undesired, unexpected events

Bad Thing;

have to make sure they are unlikely events (unlikely enough!)

Strategies:

(1) improve hardware (duplex, nplec, ecc, ...)

(2) pray

PUT

PUT(*address* , *block*)

PUT(*i* , *b*) events

Desired:

- page(*i*) = *b*, status(*i*) = good

Undesired:

Expected:

- page(*i*) not changed or status(*i*) = bad
occurs at most *k* times in a row

Unexpected:

- take your pick!


Assume: no other undesired, expected events.

especially: disk pages do not go bad on their own (decay)

Processor Model

Desired events:

- code executed correctly
- memory never lost, corrupted
- code never lost, corrupted



Undesired events:

Expected:

- fail stop
- processor halts cleanly
- memory lost
- code not lost or corrupted
- after delay, execution resumes at recovery point

Unexpected:

- what have you.

Goal: recover from failures / errors

Example: soft read errors (i.e. page is good, but GET returns STATUS=bad)

```
(st, block) = careful_get( i )
begin
  for j = 1 to (k+1) do begin
    (st, block) = get( i )
    if st = good then
      return ( st, block )
    end
  return (bad, -)
end
```


Progress

Have “eliminated” undesired, expected disk read event!


Can we do the same thing for write?

Processor failure during write:

- new value lost
- old value destroyed




Solution: **stable page**



Note:

- ignore concurrent operations:



⇒ need locking, see Chapter 7.


Careful and Stable Put

```
careful_put( i, block)
begin
    for j = 1 to (k+1) do begin
        put( i, block)
        (st, block') = careful_get( i)
        if block=block' than return
    end
    // bad block - should not get here in our model
end
```

```
stable_put( i, block)
begin
    careful_put( L( i), block)
    careful_put( R( i), block)
end
```

Recovery

- need to clean up after failures



```
recovery()  
begin  
    (stL, blockL) = careful_get( L( i) )  
    (stR, blockR) = careful_get( R( i) )  
    if stL = good and stR = bad then  
        careful_put( R( i), blockL )  
    if stL = bad and stR = good then  
        careful_put( L( i), blockR )  
    if stL = good and  
       stR = good and  
       blockL != blockR then  
        careful_put( R( i), blockL )  
end
```



stable_get

- which page (L(i) or R(i)) do we read?

```
(st, block) = stable_get( i )  
  begin  
    (st, block) = careful_get( L( i ) )  
    return ( st, block )  
  end
```

- it doesn't matter which one we read, they are the same
- might want to alternate reads for performance

What Have We Achieved?




- Atomicity
- Durability

Decay

- decay = pages go bad on their own

Reasons for Decay

- dust
- PUT (j)
- cosmic ray
- head crash
- flood
- ...



Undesired, expected events:


- set of L (or R) pages goes bad
- other pages not affected
- no other decays for at least T_d sec.

Decay Set:

- pages that can fault together, e.g.
 - Disk Drive
 - Disks in single building
 - ...

Coping with Decays


- put $L(i)$, $R(i)$ in different decay sets
- every T_d seconds: execute recovery procedure
- `stable_get` needs to cope with decay




```
(st, block) = stable_get( i )
begin
  (st, block) = careful_get( L( i ) )
  if st==good then return ( st, block)
  (st, block) = careful_get( R( i ) )
  return ( st, block)
end
```

Differences with textbook

- n-plex copies of each stable page




- no "careful" write: write n-plex copies without checking. Shotgun approach!
 - better performance
 - not as careful
 - needs version numbers



- reads needs to read **all** n-plex copies (very expensive)
- `reliable_read` does recovery: fixes bad or out-of-date pages it sees.

- ignores soft read error! (bug)




Summary

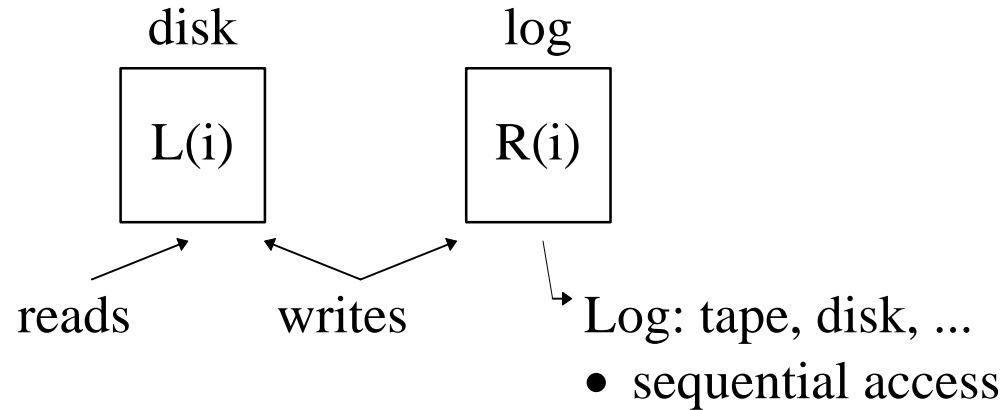
- textbook approach: efficient writes (e.g. good for log)
- stable get, put: efficient reads (Lampson, Sturgis)

Expensive?


- double, triple, ... storage
- expensive reads & writes
- recovery: read all data
- used in practice:
 - logs, mirrored disks, control blocks, ...
- optimizations/ variations:
 - be sloppy, e.g. optimistic read, Sec. 3.7.2.5
 - one copy off-line



- one copy in log



- record active page ids in log



upon recovery:

- no need to check page j
- need to check page i


Final Note

- shown how to get durability and
- atomicity of disk write

BUT NOT

- atomicity of several writes
- isolation

Example Megatron 14X disk (14 pages)



Transaction:

`stable_put(2, ...)`

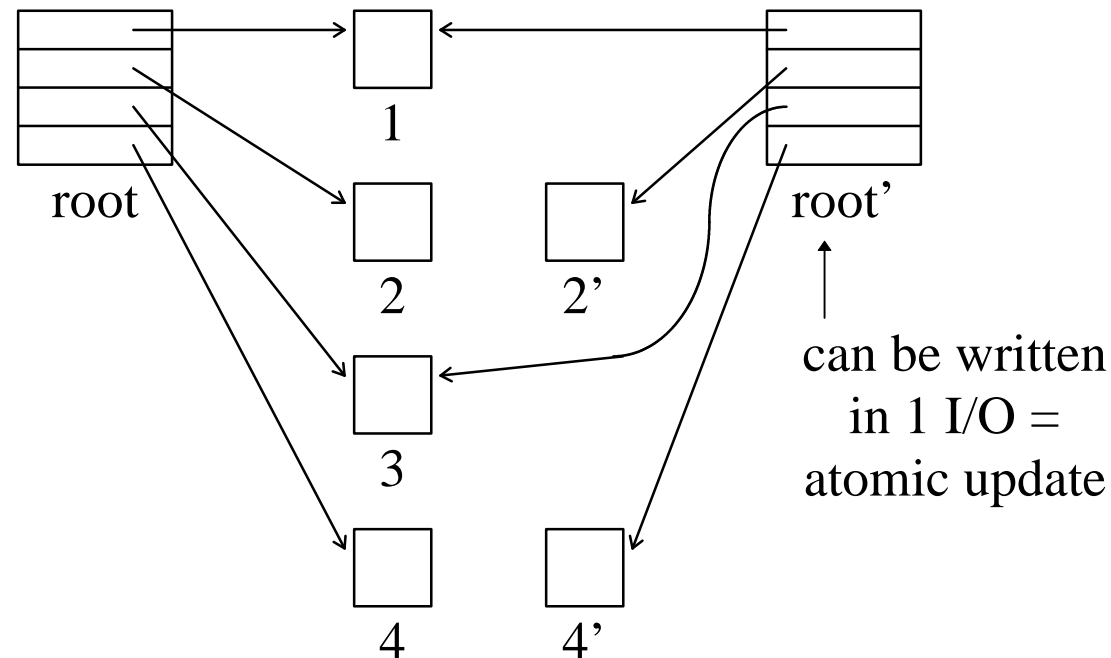
`stable_put(4, ...)`

Solutions:

1) Logging

- write to log update to page 2
- write to log update to page 4
- commit
- `stable_put(2, ...)`
- `stable_put(4, ...)`

2) shadow pages



Notes:

- can have > 2 levels
- expensive for small changes: write new version, write root, de-allocate
- better for large changes, e.g. file system
- focus of book on logging

Final Final Note

highly available processes

Read in textbook about


- checkpoint - restart
- process pairs
- persistent processes

Corrections to Example 2

Handout #10


- 3 components, MTTF = 10 years
- voter does not fail
- fail vote: system up while at least 2 up
- fail fast: system up while at least 1 up

(a) fail vote, no repairs



$$MTTF_{system} = \frac{5}{6} MTTF = 8.3 \text{ years}$$

(b) fail fast, no repairs



$$MTTF_{system} = \frac{11}{6} MTTF = 18.3 \text{ years}$$

(c) fail vote, with repairs, MTTR = 1 day = 0.003 years

“one component is down” happens with probability:

$$P_1 = (1 - \text{avail}) = \left[1 - \frac{MTTF}{MTTF + MTTR} \right] = \left[\frac{MTTR}{MTTF + MTTR} \right] \approx \left[\frac{MTTR}{MTTF} \right]$$

The system becomes unavailable when

“one specific component already down and one or both of the others fail”

The probability of this event is:


$$P_2 = P_1 \cdot P(\text{either one or both other fail})$$

$$= (1 - \text{avail}) \cdot \left[\frac{1}{MTTF} + \frac{1}{MTTF} - \frac{1}{MTTF^2} \right] \approx \frac{MTTR}{MTTF} \cdot \left[\frac{2}{MTTF} \right] = \frac{2 \cdot MTTR}{MTTF^2}$$

$$\Rightarrow P_3 = P(\text{any one component down and one or both others fail}) = 3 \cdot P_2$$

$$\Rightarrow MTTF_{\text{system}} = \frac{1}{P_3} = \frac{MTTF^2}{6 \cdot MTTR} = 5555 \text{ years}$$

(d) fail fast, with repairs, MTTR = 1 day = 0.003 years



“two components are down” happens with probability:

$$(1 - \text{avail})^2 = \left[1 - \frac{MTTF}{MTTF + MTTR} \right]^2 = \left[\frac{MTTR}{MTTF + MTTR} \right]^2 \approx \left[\frac{MTTR}{MTTF} \right]^2$$

$$\Rightarrow \text{freq}_{\text{system}} = \text{freq}_{\text{system-component}} \left[\frac{MTTR}{MTTF} \right]^2 = 3 \text{ freq}_{\text{component}} \left[\frac{MTTR}{MTTF} \right]^2$$

$$\Rightarrow MTTF_{\text{system}} = \frac{MTTF}{3} \left[\frac{MTTF}{MTTR} \right]^2 = \frac{10}{3} \left[\frac{10}{0.003} \right]^2 \text{ years} = 3.7 \cdot 10^7 \text{ years}$$